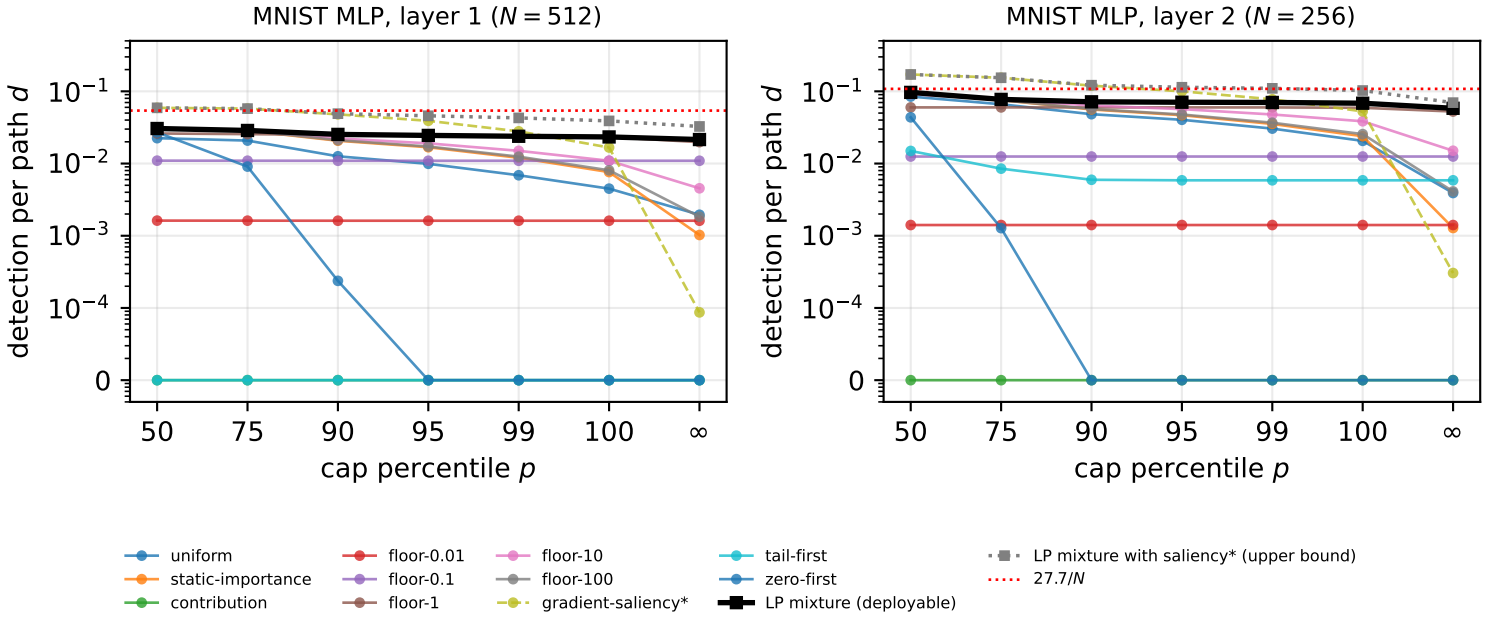
MNIST MLP, layer 1 ($N = 512$)
MNIST MLP, layer 2 ($N = 256$)
detection per path d
10^{-1}
10^{-2}
10^{-3}
10^{-4}
0
50
75
90
95
99
100
∞
cap percentile p
uniform
static-importance
contribution
floor-0.01
floor-0.1
floor-1
floor-10
floor-100
gradient-saliency*
tail-first
zero-first
LP mixture (deployable)
LP mixture with saliency* (upper bound)
27.7/N